

UAT Questionnaire v2.0 (Trimmed)

Document version: 2026-05-06 v2.0 **Source:** Trimmed from v1 (497 questions) by removing cross-role duplicates, internal implementation checks, and trivial UX repeats. **Format:** Yes / No / N/A per question, grouped by feature area, split by role.

What was cut and why

Category cut	Reason	Approx. dropped
Cross-role login / terms / logout / profile duplicates	Same flow tested 3x in v1 (ADMIN + MANAGER + USER). Now tested in full only in ADMIN; MANAGER and USER do a single smoke check that they can sign in / out and see their role enforced correctly.	~40
Implementation-only checks	Things a UAT tester cannot actually observe from the UI (audit log row was written, S3 file deleted, KB_Logs entry, JWT auto-refresh, name sync to quota record, request_id correlation). Moved to QA / dev verification.	~30
Trivial UX repeats	"Loading indicator appears", "messages are visually distinct", "alternating row colors", "tooltips show", "smooth scrolling", "page counter shows correctly" - these are visual QA, not UAT acceptance criteria. Kept one per area, dropped repeats.	~35
Sub-feature over-coverage	E.g. v1 had 7 questions for Search, 7 for Sort, 7 for Filter, 7 for Pagination. Collapsed to one combined "search + sort + filter + paginate" check per list view.	~40
Edge cases / non-functional noise	KeepAlive, mobile breakpoints, error modal styling, SweetAlert theming, browser back-button, network-error rendering. Out of scope for business UAT.	~25
Analysis-Report sub-page granularity	v1 listed every report variant (ABC, OPR, One-Page, Workload). Collapsed to "can you reach each report and generate a result".	~15
Document Extraction history granularity	v1 had 14 history-modal questions covering search / sort / order / paginate / view / delete / version. Collapsed to "open history, find, view, delete, see versions".	~20

Result: 497 -> ~256 questions. Per-role: ADMIN 120, MANAGER 70, USER 66.

v2.0.1 update (2026-05-08): Section 4 (Analysis Reports) for the User role expanded with a dedicated Workload Analysis sub-section (10 questions) covering item-row entry, configurable phase rates, calculate / totals / breakdown, validation, PDF export, save-to-DB, and offline behaviour. Original 28-question source list was de-duplicated and trimmed to 10.

The intent is that a tester can finish their role's section in one sitting (60-90 min) instead of half a day, while still covering every feature that ships to production.

ADMIN

1. Login, Terms, Session

#	Question	Yes	No
1	Can you successfully sign in with your Google account, granting Gmail / Drive / Docs / Sheets / Calendar scopes?		
2	After successful login, are you taken to the Admin dashboard with your name, email, and picture in the sidebar?		
3	If you are a new account that has not yet accepted the Terms, does a fullscreen Terms modal block the rest of the app?		
4	Is the Accept button disabled until you scroll to the bottom of the Terms?		
5	After accepting Terms once, does the modal NOT reappear on subsequent logins?		
6	If you reload the page or come back later in the same day, are you still logged in (no forced re-auth within token lifetime)?		
7	Can you log out from the sidebar, and after logout are you redirected to the login page and blocked from protected URLs?		
8	If you decline a required Google scope on consent, does the system show a clear error instead of a blank screen?		

2. Dashboard

#	Question	Yes	No
1	Does the Dashboard load with welcome card, stat cards, and quick-action shortcuts?		
2	Does the "Document Extraction" shortcut navigate to /document-extraction?		
3	Does the "Token Management" shortcut navigate to /quota?		

3. Manage Accounts

3.1 List, Search, Sort, Filter, Paginate

#	Question	Yes	No
1	Can you navigate to the Accounts page and see Activated Accounts (10 per page)?		
2	Can you toggle to "View Deactivated" and back to "View Active"?		
3	Can you search by name and email in BOTH Activated and Deactivated lists?		
4	Can you sort by Name / Email / Role / Status by clicking the column headers (and toggle asc/desc)?		
5	Can you filter by role (All / Admin / Manager / User)?		
6	Does pagination correctly trigger when the user list exceeds 10 entries (the 11th user moves to page 2)?		

3.2 Onboard, Edit, Deactivate, Reactivate

#	Question	Yes	No
1	Can you onboard a new user with Full Name, Gmail, and Role (admin / manager / user)?		
2	Does the new user immediately appear in Activated Accounts?		
3	Within ~5 minutes of onboarding, does the new user / manager / admin receive a welcome (onboarding) email at the Gmail you entered, with their assigned role and login instructions?		
4	If the welcome email does not arrive, can you confirm with the recipient that it is not in Spam / Promotions before re-onboarding?		
5	If you onboard a Gmail that previously existed (deactivated), does the system reactivate it instead of erroring?		
6	Can you click Edit on an account and successfully update full name and role?		
7	Can you deactivate an active account (with confirmation prompt) and see it move to the Deactivated list?		
8	Is the deactivated user blocked from logging in?		
9	Can you reactivate a deactivated account and confirm the user can log in again?		
10	Can you click "View Logs" on an account row and arrive at Admin Activity Logs filtered to that user?		

4. Admin Activity Logs

#	Question	Yes	No
1	Does the Admin Activity Logs page show all admin actions (onboard / update / deactivate / reactivate / accept_terms) sorted newest-first?		
2	Are entries showing timestamp, performed-by (admin name + email), action, affected user, what changed, and IP?		
3	Can you filter logs by Action Type and Time Range (e.g. last 7 / 30 days)?		
4	When you arrive from Accounts via "View Logs", are entries pre-filtered to that target user?		
5	Does pagination work (10 entries per page) and can you navigate between pages?		

5. Document Extraction (KB Upload)

5.1 Upload, Preview, Parse

#	Question	Yes	No
1	Can you click "Browse Files" and select a PDF document (and is non-PDF rejected with a clear error)?		
2	After selecting a file, are the file name and size displayed correctly, and can you clear / replace it?		
3	Does the PDF preview render all pages correctly with page numbers and is multi-page layout readable?		
4	Can you click "Parse PDF" and see a progress indicator while processing?		
5	After parsing, are chunks displayed with metadata (Page, Type, Section / Section Title)?		
6	Can you click on a chunk to highlight its location in the PDF preview, and does the PDF auto-scroll to the page?		

#	Question	Yes	No
7	Can you switch between Markdown and JSON views of the chunks, and are tables / lists / headings preserved in Markdown?		
8	Are images / image-only chunks correctly identified (type=image)?		

5.2 Edit Chunks

#	Question	Yes	No
1	Can you click Edit on a chunk to open the rich-text editor?		
2	Can you modify text and tables, save, and have the changes persist on the chunk?		
3	Can you cancel edits without saving?		

5.3 Upload to KB & Duplicates

#	Question	Yes	No
1	Does the "Upload to Knowledge Base" button appear after parsing, and does clicking it show a loading indicator?		
2	If a duplicate document is detected, does a duplicate modal appear with the existing document's details and cost-savings info?		
3	After a successful upload, does a success modal show the chunk count and filename?		
4	After a successful upload, can you immediately query the new content via the SFX Bot?		

5.4 Processing History

#	Question	Yes	No
1	Can you open the processing history modal and see all previously uploaded documents (filename, date, chunks, size, uploaded by)?		
2	Can you search, sort (date / filename / chunks / size), and paginate (10 per page) the history list?		
3	Can you view the parsed content of a historical document?		
4	Can you view a document's version history (re-uploads of the same filename)?		
5	Can you delete a document (with confirmation), and after deletion is it removed from history AND no longer returned by SFX Bot queries?		

5.5 Reset

#	Question	Yes	No
1	Can you click "Upload New PDF" to reset the preview and parsed content and process a new file?		

6. KB Analytics

#	Question	Yes	No
1	Does the KB Analytics page load with stat cards (total documents, sessions, tokens, cost) and Health Status panel?		
2	Can you toggle Document / Chat / Cost / Activity-Logs sections?		
3	Can you change the time range (e.g. 7 / 30 / 90 days) and see the metrics refresh?		

7. SFX Bot

#	Question	Yes	No
1	Can you access the SFX bot, type a message, and send it (button or Enter)?		
2	Does the bot show a thinking indicator and stream the response token-by-token?		
3	Can you click Stop to cancel a streaming response mid-flight?		
4	Does the bot respond within ~10 seconds for the first token?		
5	Does the bot retrieve relevant info from the KB and cite source documents (filename, page) in the answer?		
6	If you ask about something NOT in the KB, does the bot say so instead of hallucinating?		
7	Can you create a new chat session, switch between sessions, rename, and delete them?		
8	Does the active session persist across page refreshes?		
9	If a query exceeds your monthly quota, does a Quota Exceeded modal appear with a clear message?		
10	When the LLM service errors, does an LLM Error modal explain the failure (instead of silently breaking)?		

8. AI Assistant

8.1 Threads & Connectivity

#	Question	Yes	No
1	Can you access the AI Assistant page and see the recents / threads sidebar with prior chats?		
2	Can you start a new chat, rename, and delete a thread?		
3	Does the active thread persist across page refreshes?		
4	Is the assistant streaming responses in real time (WebSocket) and does it gracefully fall back to HTTP if WS fails?		

8.2 Clear Tasks (per Google Workspace agent)

#	Question	Yes	No
1	When you ask to create a Google Doc with sections X, Y, Z, is the doc created with those headings?		
2	When you ask to schedule a Calendar event ("tomorrow 9 AM, 30 min, with X@y.com"), is the event created with correct date / time / duration / attendees?		

#	Question	Yes	No
3	When you ask to draft an email and then send it, are recipient / subject / body correct, and does the assistant confirm before sending?		
4	When you ask to create a Sheet with tabs and headers, does the new sheet have the expected structure?		
5	When you ask the assistant to find a Drive file by name, does it correctly resolve the file (even if you paste the full Drive URL)?		
6	When you ask for a multi-step task (e.g. create event + create doc with content + attach doc link to event description), does it execute the steps in the correct order and produce the right end state?		
7	When you upload a Word template + ask to merge with a Drive data file, does the assistant produce the merged doc?		

8.3 Vague / Disambiguation

#	Question	Yes	No
1	When you give a vague request, does the assistant ask the right clarifying question before acting?		
2	When the assistant returns multiple matches ("I found 3 emails - which one?"), can you reply naturally (1, the second one, by name) to pick one?		
3	When you provide wrong / incomplete details, does the assistant detect the issue and guide you to fix it?		

8.4 Approvals (DANGEROUS / CRITICAL actions)

#	Question	Yes	No
1	When you ask for a high-impact action (send email, delete event, edit doc), does the assistant request confirmation first AND show the exact action it will take?		
2	Can you approve naturally in chat (yes, confirm) and have the workflow continue with any remaining steps?		
3	Can you reject naturally (no, cancel) and have the assistant cancel the pending action cleanly?		
4	If you reload the page mid-approval, is the pending action still there when you return?		

8.5 Files

#	Question	Yes	No
1	Can you attach a file via the paperclip (PDF / DOCX / TXT / image) and see the chip in the composer?		
2	Can you remove the file and re-attach the SAME file again (no silent no-op)?		
3	When you send a file, does the assistant successfully reference the file content?		
4	If you upload a PDF and say "use this template, find the data file in my Drive", does the assistant correctly use the upload + Drive lookup (not demand re-upload of the data file)?		

8.6 Streaming, Memory, Errors

#	Question	Yes	No
1	While the assistant works, is a step-by-step progress panel visible (in-progress / done / skipped / failed)?		
2	Does the assistant's text response stream token-by-token, and can you click Stop to cancel?		
3	If a tool fails, does the assistant explain the problem clearly and suggest a useful next step?		
4	When you continue the conversation, does the assistant correctly remember earlier context (e.g. "the doc you just created")?		
5	When you change your instruction mid-conversation, does it adapt without losing the new intent?		
6	When you ask for output formatting (e.g. "as a table" / "JSON only"), does it follow the requested format?		
7	When you ask for something the assistant cannot do (e.g. "send a fax"), does it state the limitation without hallucinating?		
8	When you exceed your monthly quota or the LLM errors, does the appropriate modal appear (Quota Exceeded / LLM Error)?		
9	Can you click "Token Usage" in the thread header and see input / output tokens, total tokens, and estimated cost for that thread?		

9. Logs & Analytics (Admin Dashboard)

#	Question	Yes	No
1	Does the Performance tab load with system metrics (avg response time, total requests) and per-agent cards (gmail / drive / docs / sheets / calendar / mapping)?		
2	Can you change the time range and see metrics refresh?		
3	Is the per-model pricing table rendered (input / cached_input / output USD per 1K tokens) with search and pagination?		
4	Can you edit the rate for a model, save, and see the change reflected in subsequent cost calculations?		
5	Can you view request logs with timestamp / user / agent / tool / success-failure / duration, and drill into a request_id to see all sub-calls?		

10. Token Quota Admin

#	Question	Yes	No
1	Is the Budget section visible at the top of the page, and can you view / update the platform-wide monthly budget?		
2	Does the user quota table list all users with current usage / monthly limit / tier / cost / reset date, with search and tier filter?		
3	Can you edit a user's tier (free / pro / enterprise / unlimited) and have the monthly_limit auto-update?		
4	Can you override a user's monthly_limit manually and edit reset_date?		
5	Can you reset a user's current_usage and current_cost_usd to zero (with audit row written)?		

#	Question	Yes	No
6	Can you deactivate and restore a user from the quota page?		
7	Is a "Top Users" panel rendered (token consumers in the selected window) with switchable time range?		
8	Can you click a user to see per-service usage breakdown (chat / agent / kb / embedding / document) with model and operation lists?		
9	Is the "Admin Actions" audit panel visible at the bottom (admin / action / target / timestamp) with pagination?		
10	When a user / manager crosses 80% of their monthly token limit, does the platform send a "near-quota" notification email to the admin(s)?		
11	When a user / manager fully exceeds (100%) their monthly token limit, does the platform send a "quota exceeded" notification email to the admin(s) with the user's identity and overage info?		
12	When the platform-wide monthly budget approaches / exceeds the configured threshold, does the admin receive a budget-alert email?		
13	If you change a user's tier or override their monthly_limit, does the next quota notification email reflect the new limit (not the old one)?		
14	If multiple admins are configured, do all of them receive the quota / budget notification emails (or only the designated recipient, per your configuration)?		

11. Profile

#	Question	Yes	No
1	Does the profile page show your full name, @username, email, role badge, member-since date, and Google avatar?		
2	Is the Token Consumption Panel visible with current usage / limit / cost / tier?		
3	Can you change the time range and toggle "Include resets"?		
4	Are SFX Bot chat turns clustered correctly (one row per conversation, not per LLM call)?		
5	Is the monthly history bar chart rendered?		

12. Cross-cutting

#	Question	Yes	No
1	When you switch pages and come back, is the previous page state preserved (KeepAlive)?		
2	Are global JavaScript errors caught and shown in an Error Modal (instead of crashing the app)?		
3	Does the app behave correctly on a narrow screen (< 1024px) - sidebar collapses, drawers toggle?		
4	After all admin testing, would you trust the assistant + KB + quota system to support real operations?		

MANAGER

Manager has access to Dashboard, AI Assistant, SFX Bot, Document Extraction (per current routing), Dynamic Mapping, Analysis Reports, and Profile. Login / Terms / Profile are smoke-tested only - depth is in ADMIN.

1. Login & Access Control (smoke)

#	Question	Yes	No
1	After the admin onboarded your manager account, did you receive a welcome / onboarding email at your Gmail (within ~5 minutes), confirming your role as manager and how to access the platform?		
2	Can you sign in with Google, accept Terms (if first time), and land on the Dashboard with your name in the sidebar?		
3	Are exactly the manager-allowed nav items visible in the sidebar (Dashboard, AI Assistant, SFX Bot, Dynamic Mapping, Analysis Report) and admin-only items hidden (Accounts, KB Analytics, Logs, Token Management, Document Extraction)?		
4	If you try to navigate by URL to an admin-only page (e.g. /accounts, /quota, /logs, /document-extraction), are you correctly redirected to /dashboard (or otherwise blocked from uploading to the KB)?		
5	Can you log out from the sidebar and be returned to the login page?		
6	Does the QuotaWidget at the top of the sidebar show your real current usage / limit (not 0 placeholder)?		

2. Dashboard

#	Question	Yes	No
1	Does the Dashboard load with welcome card and quick-action shortcuts?		
2	Do the available shortcut cards (Document Extraction, Token Management) navigate or block correctly based on your role?		

3. AI Assistant (full depth - this is the manager's primary supervisor surface)

3.1 Threads & Streaming

#	Question	Yes	No
1	Can you access the AI Assistant, see the recents sidebar, start a new chat, rename, and delete threads?		
2	Does the active thread persist across page refreshes?		
3	Are responses streamed in real time, with a step-by-step progress panel during multi-step tasks, and can you click Stop to cancel?		

3.2 Clear Tasks

#	Question	Yes	No
1	When you ask to create a Google Doc with specific sections, is the doc created with those headings?		
2	When you ask to schedule a Calendar event with date / time / attendees, is it created correctly?		
3	When you ask to draft / send a Gmail email, are recipient / subject / body correct (with confirmation before send)?		
4	When you ask to create a Sheet with tabs and headers, does the structure match?		
5	When you ask to find a Drive file by name (or paste the full Drive URL), does the assistant resolve it correctly?		
6	When you ask for a multi-step task ("create event then create doc and attach link to event description"), does it execute in the correct order?		
7	When you upload a Word template + ask to merge with a Drive data file, does the assistant produce the merged doc?		

3.3 Vague / Disambiguation / Approvals

#	Question	Yes	No
1	When you give a vague request, does the assistant ask the right clarifying question?		
2	When the assistant returns multiple matches, can you reply naturally to disambiguate?		
3	When you ask for a high-impact action (send / delete / edit), does the assistant request confirmation in chat AND show what it will do?		
4	Can you approve / reject naturally (yes / no), and does the workflow continue or stop cleanly?		
5	If you reload the page mid-approval, is the pending action still there when you return?		

3.4 Files

#	Question	Yes	No
1	Can you attach a file (paperclip), remove it, re-attach the same file (no silent no-op), and have the assistant successfully reference the file content?		

3.5 Memory, Errors, Token Usage

#	Question	Yes	No
1	Does the assistant remember earlier context (e.g. "the doc you just created") in the same thread?		
2	When you change your instruction mid-conversation, does it adapt without losing the new intent?		
3	When a tool fails, does the assistant explain clearly and suggest a useful next step?		
4	When you exceed quota or the LLM errors, does the appropriate modal appear?		
5	When you ask for something the assistant cannot do, does it say so without hallucinating?		
6	When you ask a security-sensitive question, does it refuse to expose other users' private data?		
7	Can you click "Token Usage" in the thread header and see input / output tokens and cost for that thread?		

4. SFX Bot

#	Question	Yes	No
1	Can you access SFX Bot, send a message, see the thinking indicator, and have the response stream token-by-token (with Stop available)?		
2	Does the bot retrieve relevant info from the KB and cite source documents in the answer?		
3	If you ask about something NOT in the KB, does the bot say so instead of hallucinating?		
4	Can you create / switch / rename / delete chat sessions, and does the active session persist across refreshes?		
5	If a query exceeds your monthly quota or the LLM errors, does the appropriate modal appear?		

5. Dynamic Mapping

#	Question	Yes	No
1	Can you upload a source data file (spreadsheet / CSV) and a target template?		
2	Does the AI mapping engine suggest column mappings that you can edit / approve / reject?		
3	Can you save and reload a mapping template?		
4	Can you transform the data with the mapping applied and preview the output?		

6. Analysis Reports

#	Question	Yes	No
1	Does the Analysis Report hub render with cards (ABC, One-Page, Workload) that all open without errors?		
2	Can you generate an ABC analysis from your data and view the A / B / C category breakdown?		
3	Can you view ABC Analysis History and re-open a past result?		
4	Can you generate a One-Page Report and a Workload Analysis with charts rendered correctly?		
5	Can you view your OPR History and re-open a past OPR record?		

7. Profile & Quota

#	Question	Yes	No
1	Does the profile page show your name, @username, email, role badge (manager), avatar, and Token Consumption Panel?		
2	Are SFX Bot chat turns clustered correctly (one row per conversation), and is the monthly history bar chart rendered?		
3	When you approach 80% of your monthly limit, is a warning indicator visible, and when you exceed it are AI Assistant + SFX Bot blocked?		

8. Cross-cutting (Manager)

#	Question	Yes	No
1	When you reload the page, are you still logged in (no forced re-auth)?		
2	When you switch pages and come back, is the previous page state preserved?		
3	When different managers test, does each see only their own threads / sessions / activity?		
4	After testing, would you trust this assistant to support your day-to-day operations?		

USER

Standard User has access to SFX Bot (primary landing page), Dynamic Mapping, Analysis Reports, and Profile. Other pages are blocked.

1. Login & Access Control (smoke)

#	Question	Yes	No
1	After the admin onboarded your user account, did you receive a welcome / onboarding email at your Gmail (within ~5 minutes), confirming your role as user and how to access the platform?		
2	Can you sign in with Google, accept Terms (if first time), and land on /sfx-bot?		
3	Is your role badge (user) shown on your profile page?		
4	Are exactly the user-allowed nav items visible (SFX Bot, Dynamic Mapping, Analysis Report) and all admin / manager items hidden?		
5	If you try to navigate by URL to an admin-only or manager-only page (e.g. /accounts, /dashboard, /ai-chat-new, /quota), are you correctly redirected to /sfx-bot?		
6	Can you log out from the sidebar and be returned to the login page?		
7	Does the QuotaWidget at the top of the sidebar show your real current usage / limit?		

2. SFX Bot (your primary tool - full depth)

2.1 Chat & Streaming

#	Question	Yes	No
1	Can you access the SFX bot page, type a message, and send it (button or Enter)?		
2	Does the bot show a thinking indicator and stream the response token-by-token?		
3	Can you click Stop to cancel a streaming response mid-flight?		
4	Does the bot respond within ~10 seconds for the first token?		
5	Are responses rendered as Markdown (lists, code blocks, tables, links)?		
6	Does the welcome screen show useful suggested prompts that you can click to send?		

2.2 RAG / Document Context

#	Question	Yes	No
1	Can you ask questions about uploaded documents (HR policies, SOPs, manuals) and get relevant answers?		
2	Are source documents cited in the bot's responses (filename, page)?		
3	If you ask about something NOT in the KB, does the bot say so instead of hallucinating?		
4	When you ask the same question in different wording, do you get consistent answers?		
5	When you ask for a summary of a long document, does it keep key facts and avoid omissions?		

2.3 Sessions

#	Question	Yes	No
1	Can you create a new chat session, switch between sessions, rename, and delete them?		
2	Does the active session persist across page refreshes?		
3	On a narrow screen / mobile, does the threads drawer toggle open / closed correctly?		

2.4 Privacy & Limits

#	Question	Yes	No
1	Can you ONLY see your own chat sessions (not other users')?		
2	When you ask the bot for confidential information about other users, does it refuse?		
3	If a query exceeds your monthly quota, does a Quota Exceeded modal appear with a clear message?		
4	When the LLM service errors, does an LLM Error modal explain the failure (instead of silently breaking)?		

3. Dynamic Mapping

#	Question	Yes	No
1	Can you upload a source data file (spreadsheet / CSV) and a target template?		
2	Does the AI mapping engine suggest column mappings that you can edit / approve / reject?		
3	Can you save a mapping template, load it again, transform data, and preview the output?		
4	Can you only see / load mapping templates that you (or your team) saved (not other users')?		

4. Analysis Reports

4.1 Hub

#	Question	Yes	No
1	Does the Analysis Report hub render with cards (ABC, One-Page, Workload) that all open without errors?		

4.2 ABC Analysis

#	Question	Yes	No
1	Can you generate an ABC analysis with A / B / C category breakdown, view ABC Analysis History, and re-open a past result?		

4.3 One-Page Report

#	Question	Yes	No
1	Can you generate a One-Page Report with charts rendered correctly?		
2	Can you view your OPR History and re-open a past OPR record?		

4.4 Workload Analysis

#	Question	Yes	No
1	Can you access the Workload Analysis page and see at least one item row by default?		
2	Can you add new item rows, fill in description / pallets / items per pallet, and remove rows (with at least one row always remaining)?		
3	Can you expand the warehouse rate settings, see the default phase rates loaded, and edit the rates for Inbound Checking, Put-away, Picking, and Outbound Checking?		
4	When you click Calculate Workload with valid data, does the system produce a result with totals (time, pallets, items, workers) and a per-phase breakdown?		
5	Do the totals match the inputs (total pallets equals entered pallet count, total items equals pallets times items per pallet)?		
6	If required data is missing (no item data or empty workers field), does the system stop and show a clear message?		
7	Can you export the result to PDF, and does the PDF include the summary, phase breakdown, and items breakdown?		
8	Can you save the calculation to the database when the backend is connected?		
9	Is the Save to Database button correctly disabled when no result exists or when the backend is offline?		
10	When the backend is offline, can you still calculate a workload using the loaded default rates, and is the offline state communicated clearly?		

5. Profile & Quota

#	Question	Yes	No
1	Does the profile page show your name, email, role badge (user), avatar, and Token Consumption Panel with current usage / limit / cost / tier?		
2	Are SFX Bot chat turns clustered correctly (one row per conversation), and is the monthly history bar chart rendered?		
3	Can you ONLY see your own usage data (not other users')?		
4	When you approach 80% of your monthly limit, is a warning visible, and when you exceed it is SFX Bot blocked?		

6. Cross-cutting (User)

#	Question	Yes	No
1	When you reload the page, are you still logged in?		
2	After testing, would you trust the SFX Bot + Analysis tools to support your daily work?		

Tester Sign-off

Field	Value
Tester full name	

Field	Value
Tester role (admin / manager / user)	
Test date	
Browser & OS	
Notes / blockers	
Overall verdict (Pass / Pass with notes / Fail)	
Signature	